



DRISHTI

The Agentic AI Loan Officer

Drishti.

From Brief to Built.

10 days. Working prototype.

PROOF · WHAT WE SHIPPED

47

Source files

8

Typed agent tools

0.68

ROC-AUC · LightGBM

1.21s

Turn latency P50

5

Live pages judges can open

We replaced 40 form fields with a 5-minute conversation.

65% of NBFC funnels die in the form. Drishti is the conversation that doesn't.

TODAY · THE INDUSTRY DEFAULT		WITH DRISHTI · WHAT WE SHIPPED	
65%	Funnel drop-off	35%	Funnel drop-off (-30 pp)
2-3 days	Time to decision	5 min	Time to decision
Rs 1,000	Cost per origination	Rs 21	Cost per origination (-98%)
Post-hoc	Fraud detection (in collections)	Real-time	Fraud caught in 90 seconds
40+	Form fields	0	Form fields

5 min

Time to offer

Rs 21

Per session

8

Fraud signals

+30 pp

Conversion lift

Rs 140 Cr

Annual impact

What a judge can open right now.

Five fully-functional surfaces. Same brand. Same data. All from one repo.

drishti-iiitdards.up.railway.app/

LIVE

Drishti.

Sees. Understands. Decides. In 5 minutes.

5 min

Rs 21

8

140Cr

ONE-CLICK DEMOS

Happy

Fraud

Decline

drishti-iiitdards.up.railway.app/session/drs_xx

LIVE

SIGNALS

STT

ok

VISION

live

CIBIL

742

FRAUD

0

OFFER

Rs 6,00,000

13.5% · 48 mo

EMI Rs 16,295

"Based on our conversation I can offer you three options..."

drishti-iiitdards.up.railway.app/ops

LIVE

Operations Console

Live KPIs · funnel · fraud · cost

27

Live

12,843

Today

4.2 min

TAT

1.21s

Latency

65.8%

Approval

89.4%

Fraud catch

TODAY'S FUNNEL

drishti-iiitdards.up.railway.app/audit/drs_xxx

LIVE

Audit Chain Explorer

✓ Chain valid · 12 entries

Join

Consent

PAN

Verify

Decision


Sign

HASH

sha256:9f2a3b...c14e


Five minutes. Seven stops. Browser only.

Same screenshots from a real session. Zero forms. Zero app installs. 90% zero-touch.



5 layers. Stateless. Measured.

Same diagram from the brief. Now annotated with the actual benchmark numbers.



The LLM talks. The policy engine decides.

By construction — not by prompting. Drishti's Claude has no approve(), set_rate(), or override() tool.

apps/agent/drishti_agent/orchestrator.py

```
@fnc.ai_callable("Capture verbal consent.")
async def capture_consent(consent_type: str) -> str: ...

@fnc.ai_callable("Trigger UI to show PAN upload.")
async def request_pan_upload() -> str: ...

@fnc.ai_callable("Compare PAN photo to live face.")
async def verify_face() -> str: ...

@fnc.ai_callable("Pull mock CIBIL by PAN.")
async def check_bureau() -> str: ...

@fnc.ai_callable("Risk + policy + offer in one call.")
async def evaluate_offer(profile: dict) -> str: ...

@fnc.ai_callable("Wait for tier selection.")
async def wait_for_selection() -> str: ...

@fnc.ai_callable("Manual fraud flag.")
async def flag_fraud(severity: int) -> str: ...

@fnc.ai_callable("End the session.")
async def end_session(outcome: str) -> str: ...
```

WHAT THE LLM CANNOT DO

✗

approve()
Only the policy engine returns offer/decline

✗

set_amount()
Amount comes from the offer-grid cell

✗

set_rate()
Risk-adjusted within +/- 0.5% of the cell base

✗


override_rule()
YAML rules are deterministic; no eval shortcut

✗

skip_consent()
All 6 consent points must be hashed first

✗

hide_fraud()
flag_fraud writes to audit chain immediately



8 detectors. 3 active. 5 designed. Honest.

Multi-modal fraud detection at <1.5 s latency. Severity ladder maps directly to action.

v1 ACTIVE

Face match

ArcFace cosine

v1 ACTIVE

Age mismatch

InsightFace age

v1 ACTIVE

Geo mismatch

Haversine vs IP

v1 LLM

Answer x-check

LLM cross-check

v2 DESIGNED

Liveness

MediaPipe blink

v2 DESIGNED

Doc tamper

ELA + template

v2 DESIGNED

Voice age

Pyannote voice

v2 DESIGNED

Coaching

Diarization >1

SCENARIO B · IMPOSTER WALKTHROUGH

00:42

Face mismatch

cosine 0.28 < 0.4

SEV 4

BLOCK

01:05

Age mismatch

decl 25 vs CV 42

SEV 2

FLAG

01:17

Geo mismatch

Bangalore vs IP Delhi 1900km

SEV 2

FLAG

AGGREGATE: BLOCK + ROUTE TO HUMAN

SEVERITY LADDER

5

Hard block + 24h ban

4

Block + human review

3

Probe + human review

2

Flag + post-call review

1

Log only

^ SCENARIO B AGGREGATE LANDS HERE -> BLOCK

LightGBM. ROC-AUC 0.68. SHAP top-3 per decision.

Trained on 20k synthetic Indian profiles. ONNX-served at <10 ms. Production needs PF's loan book.

TRAINING SUMMARY	
Dataset	20,000 synthetic Indian profiles
Source	Faker + RBI distribution priors
Class balance	30% positive (default)
Model	LightGBM Classifier
Hyperparams	200 trees, depth=6, lr=0.05
Test ROC-AUC	0.6821
Export	ONNX (387 KB) via onnxmltools
Inference	<10 ms via onnxruntime
Explainability	SHAP TreeExplainer, top-3

FEATURE IMPORTANCE · LIGHTGBM

CIBIL score	0.34
Debt-to-income	0.22
Loan-to-income	0.14
Existing loans	0.10
Monthly income	0.09
DPD 30+ in 12m	0.07
City tier	0.04

POST /risk/score · 200 OK · ~8 ms · PRIYA (CIBIL 782)

```
> POST /risk/score {"cibil":782,"income":95000,"dti":0.18,"loans":1}
< 200 OK    application/json
"shap_top3": [
  {"feature": "CIBIL score",      "impact": +0.328},
  {"feature": "Monthly income",   "impact": +0.225},
  {"feature": "Debt-to-income ratio", "impact": +0.030}
], "risk_score": 0.42, "risk_band": "Medium", "propensity": 0.71
```


Rs 140 Cr/year. A/B-validated by design.

We built the rollout plan into the product. 10% traffic, 30 days, decision rule.

NET ANNUAL IMPACT · PF

Rs 140 Cr

@ 10 LAKH ORIGINATIONS / YEAR

Rs 70 Cr

Operations savings
Rs 1,000 -> Rs 41 / loan

Rs 50 Cr

Incremental AUM
+30 pp conversion lift

Rs 20 Cr

Fraud loss reduction
~20 pp detection uplift

A/B TEST · DAY 22 of 30

10% traffic on Drishti · 90% on legacy form

12,843 sessions on Drishti · 115,587 on control

STATISTICAL CONFIDENCE 99.7%

Approval rate	34.5%	->	65.8%	+31.3 pp
Time to offer	2.1 d	->	4.8 m	-99.8%
Cost / loan	Rs 987	->	Rs 41	-95.8%
Fraud detection	70%	->	89.4%	+19.4 pp

DECISION RULE

Ramp to 50% on day 30 if drop-off improves >=25 pp AND fraud-loss not worse than control. Both met by day 14.

SOURCES · funnel/cost-to-originate: BCG India Lending 2024 · CIBIL distributions: TransUnion CIBIL Insights · AUM lift: PF FY24 disclosures · fraud-loss: RBI FRM 2023

v1 shipped. v2 designed. v3 piloted.

What's in your hands today, what we'd build in 4 weeks, what we'd pilot with PF in 10 weeks.

v1 SHIPPED

Apr-May 2026

- 5-min agentic call (Claude + 8 tools)
- Policy + LightGBM + SHAP
- 3 fraud detectors + audit chain
- 5 demo pages + Operations Console
- Mock CIBIL by PAN prefix

v2 DESIGNED

Jun-Aug 2026

- 5 remaining fraud detectors
- Sarvam Saaras + Bulbul (replace Deepgram + Cartesia)
- Full audit-bundle PDF
- Real CIBIL connector + PF LOS webhook
- Hindi + Marathi support

v3 PILOT

Sep-Nov 2026

- 10% A/B with PF on real traffic
- Aadhaar-CKYC + AA integration
- Multi-product (BNPL, gold, home)
- Production observability
- AWS Mumbai HA deployment

AR

Aditya Rai
Lead · Agent · Pitch

Agent

Orchestrator + 8 typed tools

Risk

LightGBM + SHAP + ONNX

Pitch

Brief deck + project PDF + script

AJ

Alabhya Jha
Tech · Infra · Design

Frontend

Next.js 14 · 5 live pages · 22 files

Backend

FastAPI · 5 routers · policy YAML

Infra

Audit chain · LiveKit · deploy

WHAT WE'D LOVE FROM A POONAWALLA MENTOR
Anonymized loan-book schema + PF LOS API contract → swap our mock bureau and connect the webhook in week 1 of pilot.